

## R Notebook

The following is your first chunk to start with. Remember, you can add chunks using the menu above (Insert -> R) or using the keyboard shortcut Ctrl+Alt+I. A good practice is to use different code chunks to answer different questions. You can delete this comment if you like.

Other useful keyboard shortcuts include Alt- for the assignment operator, and Ctrl+Shift+M for the pipe operator. You can delete these reminders if you don't want them in your report.

```
setwd("C:/Users/munis/Desktop") #Don't forget to set your working directory before you start!
```

```
library("tidyverse")
```

```
## -- Attaching packages ----- tidyverse 1.3.0 --
```

```
## v ggplot2 3.2.1      v purrr  0.3.3
## v tibble  2.1.3      v dplyr  0.8.3
## v tidyr   1.0.0      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library("tidymodels")
```

```
## Registered S3 method overwritten by 'xts':
##   method      from
##   as.zoo.xts  zoo
```

```
## -- Attaching packages ----- tidymodels 0.0.3 --
```

```
## v broom      0.5.3      v recipes  0.1.9
## v dials      0.0.4      v rsample  0.0.5
## v infer      0.5.1      v yardstick 0.0.4
## v parsnip    0.0.5
```

```
## -- Conflicts ----- tidymodels_conflicts() --
## x scales::discard() masks purrr::discard()
## x dplyr::filter()   masks stats::filter()
## x recipes::fixed()  masks stringr::fixed()
## x dplyr::lag()       masks stats::lag()
```

```

## x dials::margin()      masks ggplot2::margin()
## x yardstick::spec()    masks readr::spec()
## x recipes::step()      masks stats::step()
## x recipes::yj_trans()  masks scales::yj_trans()

library("plotly")

##
## Attaching package: 'plotly'

## The following object is masked from 'package:ggplot2':
##
##   last_plot

## The following object is masked from 'package:stats':
##
##   filter

## The following object is masked from 'package:graphics':
##
##   layout

library("skimr")

dfc <-
  read_csv("C:/Users/munis/Desktop/assignment3Carvana.csv")

## Parsed with column specification:
## cols(
##   Auction = col_character(),
##   Age = col_double(),
##   Make = col_character(),
##   Color = col_character(),
##   WheelType = col_character(),
##   Odo = col_double(),
##   Size = col_character(),
##   MMRAuction = col_double(),
##   MMRAretail = col_double(),
##   BadBuy = col_double()
## )

dfc

## # A tibble: 10,061 x 10
##   Auction Age Make Color WheelType Odo Size MMRAuction MMRAretail
##   <chr> <dbl> <chr> <chr> <chr> <dbl> <chr> <dbl> <dbl>
##   <dbl>
## 1 MANHEIM 4 CHRY~ GOLD Covers 45224 CROS~ 0 0
1
## 2 ADESA 6 PONT~ SILV~ Covers 79905 LARGE 3912 4725
1

```

```
## 3 MANHEIM      6 SATU~ WHITE Covers      81116 MEDI~      2667      3380
0
## 4 ADESA        3 CHEV~ WHITE NULL      82651 COMP~      5194      8614
1
## 5 MANHEIM      6 CHRY~ BEIGE Alloy      69050 MEDI~      3265      6441
1
## 6 OTHER        5 CHEV~ RED    Alloy      54718 MEDI~      6921      7975
1
## 7 MANHEIM      5 FORD  SILV~ Alloy      75757 COMP~      3475      6958
1
## 8 ADESA        4 FORD  BLUE  Covers      65726 VAN      3889      4700
0
## 9 OTHER        5 CHEV~ GOLD  Covers      89365 VAN      6131      9793
1
## 10 ADESA       3 CHEV~ WHITE Covers      71794 VAN      6394      7406
0
## # ... with 10,051 more rows
```

**head**(dfc)

```
## # A tibble: 6 x 10
##   Auction Age Make   Color WheelType Odo Size MMRAauction MMRAretail
BadBuy
##   <chr>   <dbl> <chr>   <chr> <chr>   <dbl> <chr>   <dbl>   <dbl>
<dbl>
## 1 MANHEIM      4 CHRYS~ GOLD  Covers      45224 CROS~      0      0
1
## 2 ADESA        6 PONTI~ SILV~ Covers      79905 LARGE      3912      4725
1
## 3 MANHEIM      6 SATURN WHITE Covers      81116 MEDI~      2667      3380
0
## 4 ADESA        3 CHEVR~ WHITE NULL      82651 COMP~      5194      8614
1
## 5 MANHEIM      6 CHRYS~ BEIGE Alloy      69050 MEDI~      3265      6441
1
## 6 OTHER        5 CHEVR~ RED    Alloy      54718 MEDI~      6921      7975
1
```

**nrow**(dfc)

```
## [1] 10061
```

**skim**(dfc)

*Data summary*

Name	dfc
Number of rows	10061
Number of columns	10

---

Column type frequency:

character                5  
numeric                 5





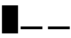
---

Group variables         None

### Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
Auction	0	1	5	7	0	3	0
Make	0	1	3	10	0	30	0
Color	0	1	3	8	0	17	0
WheelType	0	1	4	7	0	4	0
Size	0	1	3	10	0	12	0

### Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
Age	0	1	4.50	1.77	1	3	4	6	9	
Odo	0	1	72903.87	14498.87	94	634	749	836	1157	
MMRAuction	0	1	5812.38	2578.85	0	387	558	745	3572	
MMRAetail	0	1	8171.51	3257.19	0	587	805	103	3908	
BadBuy	0	1	0.50	0.50	0	0	0	1	1	


```
set.seed(52156)
```

```
#Split the data into two: 80% for the training set, and 20% for the test set
```


```
dfcTrain <- dfc %>% sample_frac(0.65)
```

```
dfcTest <- dplyr::setdiff(dfc, dfcTrain)
```


```
ggplot(data=dfcTrain)+ geom_boxplot(aes(x=factor(BadBuy),y=MMRAuction),fill="lightblue")
```



```
ggplotly( ggplot(data=dfcTrain)+ geom_boxplot(aes(x=factor(BadBuy),y=MMRAuction),fill="lightblue"))  
  
ggplot(data=dfcTrain)+ geom_boxplot(aes(x=factor(BadBuy),y=Age),fill="lightblue")
```



```
ggplotly(ggplot(data=dfcTrain)+ geom_boxplot(aes(x=factor(BadBuy),y=Age),fill="lightblue"))  
  
ggplot(data=dfcTrain)+ geom_boxplot(aes(x=factor(BadBuy),y=Odo),fill="lightblue")
```



```
ggplotly( ggplot(data=dfcTrain)+ geom_boxplot(aes(x=factor(BadBuy),y=Odo),fill="lightblue"))
```

```
library(car)
```

```
## Loading required package: carData
```

```
## Registered S3 methods overwritten by 'car':
```

```
##   method                      from
##   influence.merMod             lme4
##   cooks.distance.influence.merMod lme4
##   dfbeta.influence.merMod       lme4
##   dfbetas.influence.merMod      lme4
```

```
##
```

```
## Attaching package: 'car'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##   recode
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##   some
```

```
#library(modelr)
```

```
library(caret)
```

```
## Loading required package: lattice

##
## Attaching package: 'caret'

## The following objects are masked from 'package:yardstick':
##
##   precision, recall

## The following object is masked from 'package:purrr':
##
##   lift
```

```
dfcTrain %>%
  group_by(Size) %>%
  filter(BadBuy==1)%>%
  tally(name='Lemons') %>%
  mutate(pct=100*(Lemons)/sum(Lemons)) %>%
  arrange(desc(Lemons))
```

```
## # A tibble: 12 x 3
##   Size      Lemons  pct
##   <chr>    <int> <dbl>
## 1 MEDIUM      1298 39.8
## 2 COMPACT       448 13.7
## 3 MEDIUMSUV    412 12.6
## 4 LARGE        284  8.70
## 5 VAN          269  8.24
## 6 LARGETRUCK   126  3.86
## 7 SMALLSUV     112  3.43
## 8 LARGESUV      76  2.33
## 9 SPECIALTY     68  2.08
## 10 CROSSOVER    66  2.02
## 11 SPORTS       58  1.78
## 12 SMALLTRUCK   47  1.44
```

```
dfcTrain %>%
  group_by(Size) %>%
  filter(BadBuy==0)%>%
  tally(name='Good_cars') %>%
  mutate(pct=100*(Good_cars)/sum(Good_cars)) %>%
  arrange(desc(Good_cars))
```

```
## # A tibble: 12 x 3
##   Size      Good_cars  pct
##   <chr>    <int> <dbl>
## 1 MEDIUM      1384 42.2
## 2 LARGE        423 12.9
## 3 MEDIUMSUV    348 10.6
## 4 COMPACT      309  9.43
## 5 VAN          250  7.63
```



```
## 6 LARGETRUCK      156  4.76
## 7 SMALLSUV        97   2.96
## 8 CROSSOVER       88   2.69
## 9 SPECIALTY       79   2.41
## 10 LARGESUV       53   1.62
## 11 SPORTS         46   1.40
## 12 SMALLTRUCK     43   1.31
```

Q3a)b)

```
resultsLPM1<-
  lm(BadBuy~ ., data=dfcTrain)
#detach('package:modelr', unload=TRUE)
resultsLPM1

##
## Call:
## lm(formula = BadBuy ~ ., data = dfcTrain)
##
## Coefficients:
##      (Intercept)      AuctionMANHEIM      AuctionOTHER      Age
##      -1.996e-01      4.065e-02      2.287e-02      5.154e-02
##      MakeBUICK      MakeCADILLAC      MakeCHEVROLET      MakeCHRYSLER
##      2.392e-01      2.664e-01      1.861e-01      2.944e-01
##      MakeDODGE      MakeFORD      MakeGMC      MakeHONDA
##      2.384e-01      2.620e-01      1.398e-01      1.114e-01
##      MakeHYUNDAI      MakeINFINITI      MakeISUZU      MakeJEEP
##      2.099e-01      3.671e-01      1.764e-01      2.537e-01
##      MakeKIA      MakeLEXUS      MakeLINCOLN      MakeMAZDA
##      2.190e-01      8.805e-01      3.712e-01      2.567e-01
##      MakeMERCURY      MakeMINI      MakeMITSUBISHI      MakeNISSAN
##      2.980e-01      3.301e-01      1.179e-01      2.310e-01
##      MakeOLDSMOBILE      MakePONTIAC      MakeSATURN      MakeSCION
##      3.261e-01      2.181e-01      2.800e-01      1.091e-01
##      MakeSUBARU      MakeSUZUKI      MakeTOYOTA      MakeVOLKSWAGEN
##      2.432e-01      3.696e-01      1.638e-01      2.630e-01
##      MakeVOLVO      ColorBLACK      ColorBLUE      ColorBROWN
##      -1.809e-01      2.220e-02      1.890e-02      1.819e-02
##      ColorGOLD      ColorGREEN      ColorGREY      ColorMAROON
##      5.438e-02      2.264e-02      3.804e-02      7.248e-02
##      ColorNOTAVAIL      ColorNULL      ColorORANGE      ColorOTHER
##      -4.753e-02      -1.179e-01      4.598e-02      -1.388e-01
##      ColorPURPLE      ColorRED      ColorSILVER      ColorWHITE
##      1.955e-02      6.169e-02      4.814e-02      6.047e-02
##      ColorYELLOW      WheelTypeCovers      WheelTypeNULL      WheelTypeSpecial
##      -6.072e-02      -3.534e-02      5.096e-01      -8.805e-03
##      Odo      SizeCROSSOVER      SizeLARGE      SizeLARGESUV
##      2.888e-06      -1.783e-01      -1.475e-01      -1.379e-01
##      SizeLARGETRUCK      SizeMEDIUM      SizeMEDIUMSUV      SizeSMALLSUV
##      -1.916e-01      -9.926e-02      -9.874e-02      -1.333e-01
```

```
##      SizeSMALLTRUCK      SizeSPECIALTY      SizeSPORTS      SizeVAN
##      -1.449e-01      -7.220e-02      -1.081e-01      -1.136e-01
##      MMRAuction      MMRAretail
##      1.595e-06      -1.126e-06
```

```
#dfcTrain
```

```
performance<-metric_set(rmse,mae)
```

```
result<-dfcTest%>%
```

```
  mutate('predictedVal'= predict(resultsLPM1, dfcTest))
```

```
result
```

```
## # A tibble: 3,521 x 11
```

```
##   Auction Age Make Color WheelType Odo Size MMRAuction MMRAretail
##   <chr>   <dbl> <chr> <chr> <chr>   <dbl> <chr>      <dbl>      <dbl>
##   <dbl>
```

```
## 1 MANHEIM 6 SATU~ WHITE Covers 81116 MEDI~ 2667 3380
## 2 OTHER 5 CHEV~ RED Alloy 54718 MEDI~ 6921 7975
## 3 OTHER 5 CHEV~ GOLD Covers 89365 VAN 6131 9793
## 4 ADESA 3 CHEV~ WHITE Covers 71794 VAN 6394 7406
## 5 OTHER 3 CHEV~ WHITE NULL 67229 COMP~ 5785 9834
## 6 MANHEIM 3 DODGE GOLD Covers 71079 MEDI~ 4297 5141
## 7 MANHEIM 6 OLDS~ SILV~ Alloy 71235 MEDI~ 3325 4091
## 8 MANHEIM 8 PONT~ SILV~ Alloy 90325 MEDI~ 2150 4937
## 9 MANHEIM 6 PONT~ GREEN Alloy 96893 MEDI~ 4059 4884
## 10 OTHER 2 DODGE BLUE Covers 45151 MEDI~ 7982 9121
```

```
## # ... with 3,511 more rows, and 1 more variable: predictedVal <dbl>
```

```
Model<-performance(result,truth=BadBuy,estimate=predictedVal)
```

```
Model
```

```
## # A tibble: 2 x 3
```

```
##   .metric .estimator .estimate
```

```
##   <chr>   <chr>      <dbl>
```

```
## 1 rmse    standard    0.453
```

```
## 2 mae     standard    0.415
```

```
result2<-dfcTrain%>%
```

```
  mutate('predictedVal'= predict(resultsLPM1, dfcTrain))
```

```
Model2<-performance(result2,truth=BadBuy,estimate=predictedVal)
```

```
Model2
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      0.448
## 2 mae     standard      0.410

colsToFactor<- c('BadBuy')
dfcTest<-dfcTest%>%
  mutate_at(colsToFactor, ~factor(.))

colsToFactor<- c('BadBuy')
dfcTrain<-dfcTrain%>%
  mutate_at(colsToFactor, ~factor(.))

str(dfcTrain)

## Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame': 6540 obs. of  10
## variables:
## $ Auction      : chr  "MANHEIM" "MANHEIM" "MANHEIM" "ADESA" ...
## $ Age          : num  4 5 2 4 5 4 2 7 6 8 ...
## $ Make         : chr  "FORD" "MINI" "CHEVROLET" "NISSAN" ...
## $ Color        : chr  "SILVER" "BLUE" "SILVER" "BLUE" ...
## $ WheelType    : chr  "NULL" "Alloy" "Covers" "NULL" ...
## $ Odo          : num  77591 80013 75493 84827 57388 ...
## $ Size         : chr  "LARGETRUCK" "COMPACT" "LARGE" "MEDIUM" ...
## $ MMRAuction: num  9774 11040 9707 6073 5574 ...
## $ MMRAretail : num  14506 12423 13975 9791 8984 ...
## $ BadBuy       : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 1 2 1 1 ...
```

Q3)c)

```
#colsToFactor<- c('BadBuy')
#resultsLPM<-resultsLPM%>%
#  mutate_at(colsToFactor, ~factor(.))

#resultsLPM1<-
#  lm(BadBuy~ ., data=dfcTrain)
#result<-
resultsLPM<-resultsLPM1%>%

predict(dfcTest, type='response') %>%      #=> Use the option type='response'
for probabilities
  bind_cols(dfcTest, predictedProb=.)%>%
  mutate(predictedClass = as.factor(ifelse(predictedProb>0.5, 1, 0)))

resultsLPM%>%
  group_by(predictedClass) %>%
  tally %>%
```

```

#group_by(school_number) %>%
mutate(pct=(100*n)/sum(n))#%>%

## # A tibble: 2 x 3
##   predictedClass     n    pct
##   <fct>          <int> <dbl>
## 1 0              2117  60.1
## 2 1              1404  39.9

#resultsLPM

resultsLPM%>%
  conf_mat(truth=BadBuy, estimate=predictedClass)

##           Truth
## Prediction    0    1
##           0 1374  743
##           1  408  996

newData <- tibble(Auction="ADESA",      Age=1,      Make="HONDA",      Color="SI
LVER",
WheelType="Covers", Odo=10000,  Size="LARGE",
MMRAuction=8000,  MMRAretail=10000
)
predict(resultsLPM1, newData)

##           1
## -0.1410712

dfc

## # A tibble: 10,061 x 10
##   Auction  Age Make  Color WheelType  Odo Size  MMRAuction MMRAretail
BadBuy
##   <chr>   <dbl> <chr> <chr> <chr>      <dbl> <chr>      <dbl>      <dbl>
<dbl>
## 1 MANHEIM      4 CHRY~ GOLD  Covers    45224 CROS~          0          0
1
## 2 ADESA        6 PONT~ SILV~ Covers    79905 LARGE    3912    4725
1
## 3 MANHEIM      6 SATU~ WHITE Covers    81116 MEDI~    2667    3380
0
## 4 ADESA        3 CHEV~ WHITE NULL      82651 COMP~    5194    8614
1
## 5 MANHEIM      6 CHRY~ BEIGE Alloy     69050 MEDI~    3265    6441
1
## 6 OTHER        5 CHEV~ RED   Alloy     54718 MEDI~    6921    7975
1
## 7 MANHEIM      5 FORD  SILV~ Alloy     75757 COMP~    3475    6958
1
## 8 ADESA        4 FORD  BLUE  Covers    65726 VAN     3889    4700

```

```

0
## 9 OTHER          5 CHEV~ GOLD  Covers    89365 VAN          6131          9793
1
## 10 ADESA         3 CHEV~ WHITE Covers    71794 VAN          6394          7406
0
## # ... with 10,051 more rows

dfc %>%
  group_by((Color)) %>%
  tally %>%
  #group_by(school_number) %>%
  mutate(pct=(100*n)/sum(n))

## # A tibble: 17 x 3
##   `(Color)`      n    pct
##   <chr>      <int> <dbl>
## 1 BEIGE        237  2.36
## 2 BLACK       1013 10.1
## 3 BLUE        1386 13.8
## 4 BROWN         65  0.646
## 5 GOLD         766  7.61
## 6 GREEN        442  4.39
## 7 GREY       1054 10.5
## 8 MAROON        281  2.79
## 9 NOTAVAIL       26  0.258
## 10 NULL          2  0.0199
## 11 ORANGE        43  0.427
## 12 OTHER         37  0.368
## 13 PURPLE        57  0.567
## 14 RED          881  8.76
## 15 SILVER       2081 20.7
## 16 WHITE       1653 16.4
## 17 YELLOW        37  0.368

dfc%>%
  filter(Color=="NULL" | Color=="NOTAVAIL")

## # A tibble: 28 x 10
##   Auction Age Make Color WheelType Odo Size MMRAauction MMRAretail
BadBuy
##   <chr>   <dbl> <chr> <chr> <chr>      <dbl> <chr>      <dbl>      <dbl>
<dbl>
## 1 OTHER      2 DODGE NOTA~ NULL      50104 MEDI~      7533      10574
1
## 2 MANHEIM    6 FORD  NOTA~ Alloy    86745 MEDI~      3838      4645
1
## 3 OTHER      9 FORD  NOTA~ NULL    93287 SPOR~      2827      5871
1
## 4 OTHER      1 CHRY~ NOTA~ Covers  50633 MEDI~      8252      9412
0
## 5 OTHER      5 JEEP  NOTA~ NULL    77394 SMAL~      5485      6424

```

```

1
## 6 OTHER      5 FORD  NOTA~ Alloy      71145 MEDI~      6078      7064
1
## 7 ADESA      7 MERC~ NULL  NULL      85546 MEDI~      3779      4581
1
## 8 OTHER      5 JEEP  NOTA~ Alloy      85385 MEDI~      6682      7717
1
## 9 MANHEIM    8 PONT~ NOTA~ Alloy      73430 SPOR~      4650      7935
1
## 10 ADESA     3 PONT~ NOTA~ NULL      50486 MEDI~      9763      11044
1
## # ... with 18 more rows

```

Q4)a)

```

dfc <- dfc %>% mutate(Color = if_else(Color == "NULL", "NOTAVAIL", Color))

dfc %>%
  group_by((Make)) %>%
  tally()

## # A tibble: 30 x 2
##   `(Make)`      n
##   <chr>      <int>
## 1 ACURA         9
## 2 BUICK        103
## 3 CADILLAC        3
## 4 CHEVROLET    2121
## 5 CHRYSLER    1217
## 6 DODGE        1653
## 7 FORD         1764
## 8 GMC           85
## 9 HONDA         77
## 10 HYUNDAI     239
## # ... with 20 more rows

new<-c("ACURA","VOLVO","CADILLAC","MINI","SUBARU","LEXUS")

dfc<-dfc%>%
  mutate(Make=if_else(Make %in% unlist(new),"OTHER",Make ))

dfc<-dfc%>%
  mutate(BadBuy=as.factor(BadBuy))

set.seed(52156)

dfcTrain<-dfc%>%sample_frac((0.65))

dfcTest<-dplyr::setdiff(dfc,dfcTrain)

```

Q4)d)

```
library(e1071)
library(caret)
resultsCaret<-
  train(BadBuy ~ ., family='binomial', data=dfcTrain, method='glm') %>%
  predict(dfcTest, type='raw')%>%
  bind_cols(dfcTest, predictedClass=.)

resultsCaret%>%
  xtabs(~predictedClass+BadBuy, .)%>%
  confusionMatrix(positive='1')

## Confusion Matrix and Statistics
##
##               BadBuy
## predictedClass    0    1
##               0 1341  721
##               1  441 1018
##
##               Accuracy : 0.67
##               95% CI : (0.6542, 0.6855)
##               No Information Rate : 0.5061
##               P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.3386
##
##  Mcnemar's Test P-Value : 2.731e-16
##
##               Sensitivity : 0.5854
##               Specificity : 0.7525
##               Pos Pred Value : 0.6977
##               Neg Pred Value : 0.6503
##               Prevalence : 0.4939
##               Detection Rate : 0.2891
##               Detection Prevalence : 0.4144
##               Balanced Accuracy : 0.6690
##
##               'Positive' Class : 1
##
```

Q4)b)c)

```
summary(dfc)

##      Auction      Age      Make      Color
## Length:10061   Min.   :1.000 Length:10061 Length:10061
## Class :character 1st Qu.:3.000 Class :character Class :character
## Mode  :character Median :4.000 Mode  :character Mode  :character
##                Mean   :4.505
```

```
##          3rd Qu.:6.000
##          Max.    :9.000
##   WheelType      Odo      Size      MMRAuction
## Length:10061    Min.   : 9446   Length:10061    Min.   : 0
## Class :character 1st Qu.: 63488   Class :character 1st Qu.: 3877
## Mode  :character Median : 74942   Mode  :character Median : 5588
##          Mean    : 72904          Mean    : 5812
##          3rd Qu.: 83663          3rd Qu.: 7450
##          Max.    :115717         Max.    :35722
##   MMRAretail    BadBuy
## Min.   : 0      0:5058
## 1st Qu.: 5872   1:5003
## Median : 8052
## Mean   : 8172
## 3rd Qu.:10315
## Max.   :39080
```

```
model <- glm(BadBuy ~ ., family='binomial', data=dfcTrain)
summary(model)
```

```
##
## Call:
## glm(formula = BadBuy ~ ., family = "binomial", data = dfcTrain)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0725  -0.9782  -0.4717   1.0946   2.1705
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.472e+00  4.513e-01  -5.478 4.30e-08 ***
## AuctionMANHEIM 1.735e-01  7.493e-02   2.316 0.020579 *
## AuctionOTHER   9.519e-02  9.037e-02   1.053 0.292217
## Age            2.785e-01  2.887e-02   9.647 < 2e-16 ***
## MakeCHEVROLET -2.774e-01  2.895e-01  -0.958 0.337982
## MakeCHRYSLER   2.527e-01  3.011e-01   0.839 0.401419
## MakeDODGE      -2.483e-02  2.966e-01  -0.084 0.933287
## MakeFORD        1.020e-01  2.945e-01   0.346 0.729155
## MakeGMC         -5.054e-01  4.193e-01  -1.205 0.228054
## MakeHONDA       -6.530e-01  4.317e-01  -1.512 0.130433
## MakeHYUNDAI    -1.623e-01  3.381e-01  -0.480 0.631275
## MakeINFINITI    3.727e-01  1.280e+00   0.291 0.771007
## MakeISUZU      -3.227e-01  7.887e-01  -0.409 0.682408
## MakeJEEP        3.121e-02  3.496e-01   0.089 0.928850
## MakeKIA         -9.342e-02  3.281e-01  -0.285 0.775823
## MakeLINCOLN     6.866e-01  7.410e-01   0.927 0.354146
## MakeMAZDA       3.015e-02  3.530e-01   0.085 0.931925
## MakeMERCURY     2.670e-01  3.632e-01   0.735 0.462313
## MakeMITSUBISHI -6.722e-01  3.692e-01  -1.821 0.068664 .
## MakeNISSAN      -7.824e-02  3.213e-01  -0.243 0.807645
```



```

## MakeOLDSMOBILE      4.725e-01  5.397e-01   0.875  0.381344
## MakeOTHER            3.109e-01  6.256e-01   0.497  0.619240
## MakePONTIAC          -1.156e-01  3.039e-01  -0.380  0.703748
## MakeSATURN           2.040e-01  3.293e-01   0.620  0.535513
## MakeSCION            -6.429e-01  7.485e-01  -0.859  0.390426
## MakeSUZUKI           6.756e-01  3.578e-01   1.888  0.058974 .
## MakeTOYOTA           -4.609e-01  3.718e-01  -1.240  0.215081
## MakeVOLKSWAGEN       3.278e-02  6.815e-01   0.048  0.961638
## ColorBLACK           1.502e-01  2.157e-01   0.696  0.486312
## ColorBLUE            1.197e-01  2.103e-01   0.569  0.569124
## ColorBROWN           1.348e-01  3.891e-01   0.346  0.729074
## ColorGOLD            3.066e-01  2.201e-01   1.393  0.163652
## ColorGREEN           1.723e-01  2.369e-01   0.727  0.466976
## ColorGREY            2.307e-01  2.139e-01   1.078  0.280903
## ColorMAROON           4.114e-01  2.596e-01   1.585  0.112963
## ColorNOTAVAIL        -2.898e-01  7.521e-01  -0.385  0.700011
## ColorORANGE          2.922e-01  4.655e-01   0.628  0.530251
## ColorOTHER           -1.168e+00  6.442e-01  -1.812  0.069933 .
## ColorPURPLE           1.899e-01  4.250e-01   0.447  0.655029
## ColorRED              3.374e-01  2.177e-01   1.550  0.121257
## ColorSILVER           2.850e-01  2.057e-01   1.386  0.165860
## ColorWHITE           3.409e-01  2.083e-01   1.636  0.101745
## ColorYELLOW          -2.904e-01  4.947e-01  -0.587  0.557141
## WheelTypeCovers      -1.082e-01  6.698e-02  -1.615  0.106304
## WheelTypeNULL        3.489e+00  1.727e-01  20.202  < 2e-16 ***
## WheelTypeSpecial     -5.363e-02  2.663e-01  -0.201  0.840390
## Odo                   1.484e-05  2.184e-06   6.796  1.08e-11 ***
## SizeCROSSOVER        -9.331e-01  2.220e-01  -4.203  2.63e-05 ***
## SizeLARGE            -7.613e-01  1.319e-01  -5.770  7.91e-09 ***
## SizeLARGESUV         -7.972e-01  2.454e-01  -3.249  0.001157 **
## SizeLARGETRUCK       -1.013e+00  1.827e-01  -5.547  2.90e-08 ***
## SizeMEDIUM           -5.260e-01  1.015e-01  -5.181  2.21e-07 ***
## SizeMEDIUMSUV        -5.453e-01  1.425e-01  -3.826  0.000130 ***
## SizeSMALLSUV         -6.989e-01  2.079e-01  -3.361  0.000776 ***
## SizeSMALLTRUCK       -7.329e-01  2.520e-01  -2.908  0.003632 **
## SizeSPECIALTY        -4.271e-01  2.274e-01  -1.878  0.060352 .
## SizeSPORTS           -5.701e-01  2.545e-01  -2.240  0.025066 *
## SizeVAN              -5.982e-01  1.362e-01  -4.394  1.11e-05 ***
## MMRAuction           2.895e-05  3.634e-05   0.797  0.425670
## MMRAretail           -8.784e-06  2.241e-05  -0.392  0.695044
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 9066.3  on 6539  degrees of freedom
## Residual deviance: 7528.1  on 6480  degrees of freedom
## AIC: 7648.1
##
## Number of Fisher Scoring iterations: 5

```

```
exp(coef(model))
```

```
##      (Intercept) AuctionMANHEIM AuctionOTHER      Age
##      0.08437858      1.18946691      1.09986387      1.32116202
##      MakeCHEVROLET MakeCHRYSLER      MakeDODGE      MakeFORD
##      0.75778178      1.28744383      0.97547613      1.10735004
##      MakeGMC      MakeHONDA      MakeHYUNDAI      MakeINFINITI
##      0.60323457      0.52049737      0.85021505      1.45161848
##      MakeISUZU      MakeJEEP      MakeKIA      MakeLINCOLN
##      0.72416820      1.03170472      0.91080706      1.98696169
##      MakeMAZDA      MakeMERCURY      MakeMITSUBISHI      MakeNISSAN
##      1.03061116      1.30598879      0.51060714      0.92474703
##      MakeOLDSMOBILE      MakeOTHER      MakePONTIAC      MakeSATURN
##      1.60397687      1.36460391      0.89085880      1.22630987
##      MakeSCION      MakeSUZUKI      MakeTOYOTA      MakeVOLKSWAGEN
##      0.52579005      1.96521483      0.63068658      1.03332045
##      ColorBLACK      ColorBLUE      ColorBROWN      ColorGOLD
##      1.16204597      1.12720518      1.14428128      1.35881972
##      ColorGREEN      ColorGREY      ColorMAROON      ColorNOTAVAIL
##      1.18807424      1.25947392      1.50895187      0.74842483
##      ColorORANGE      ColorOTHER      ColorPURPLE      ColorRED
##      1.33931201      0.31112534      1.20909575      1.40123691
##      ColorSILVER      ColorWHITE      ColorYELLOW      WheelTypeCovers
##      1.32980630      1.40616624      0.74793801      0.89747194
##      WheelTypeNULL      WheelTypeSpecial      Odo      SizeCROSSOVER
##      32.73825555      0.94777937      1.00001484      0.39333606
##      SizeLARGE      SizeLARGESUV      SizeLARGETRUCK      SizeMEDIUM
##      0.46705813      0.45057313      0.36300261      0.59094577
##      SizeMEDIUMSUV      SizeSMALLSUV      SizeSMALLTRUCK      SizeSPECIALTY
##      0.57963900      0.49710890      0.48052501      0.65237868
##      SizeSPORTS      SizeVAN      MMRAauction      MMRAretail
##      0.56544253      0.54978912      1.00002895      0.99999122
```

Q4)d)

```
result <-
  glm(BadBuy ~ ., family='binomial', data=dfcTrain) %>%
  predict(dfctest, type='response') %>%
  bind_cols(dfctest, predictedProb=.) %>%
  mutate(predictedClass = as.factor(ifelse(predictedProb>0.5, 1, 0)))

result %>%
  group_by(predictedClass) %>%
  tally %>%
  #group_by(school_number) %>%
  mutate(pct=(100*n)/sum(n))

## # A tibble: 2 x 3
##   predictedClass     n    pct
##   <fct>           <int> <dbl>
```

```
## 1 0          2062  58.6
## 2 1          1459  41.4

result%>%
  conf_mat(truth=BadBuy, estimate=predictedClass)

##           Truth
## Prediction    0    1
##           0 1341  721
##           1  441 1018

result

## # A tibble: 3,521 x 12
##   Auction Age Make Color WheelType Odo Size MMRAuction MMRAretail
BadBuy
##   <chr>   <dbl> <chr> <chr> <chr>      <dbl> <chr>      <dbl>      <dbl>
<fct>
## 1 MANHEIM     6 SATU~ WHITE Covers   81116 MEDI~    2667    3380
0
## 2 OTHER      5 CHEV~ RED Alloy    54718 MEDI~    6921    7975
1
## 3 OTHER      5 CHEV~ GOLD Covers   89365 VAN     6131    9793
1
## 4 ADESA      3 CHEV~ WHITE Covers   71794 VAN     6394    7406
0
## 5 OTHER      3 CHEV~ WHITE NULL    67229 COMP~    5785    9834
1
## 6 MANHEIM     3 DODGE GOLD Covers   71079 MEDI~    4297    5141
1
## 7 MANHEIM     6 OLDS~ SILV~ Alloy    71235 MEDI~    3325    4091
1
## 8 MANHEIM     8 PONT~ SILV~ Alloy    90325 MEDI~    2150    4937
1
## 9 MANHEIM     6 PONT~ GREEN Alloy    96893 MEDI~    4059    4884
1
## 10 OTHER      2 DODGE BLUE Covers   45151 MEDI~    7982    9121
1
## # ... with 3,511 more rows, and 2 more variables: predictedProb <dbl>,
## #   predictedClass <fct>
```

Q4)e)

```
newData2 <- tibble(Auction="ADESA", Age=1, Make="HONDA", Color="SILVER",
  WheelType="Covers", Odo=10000, Size="LARGE",
  MMRAuction=8000, MMRAretail=10000
)
predict(model, newData2)
```

```
##          1
## -3.139145
```

Q5)a)

```
set.seed(123)

resultsCaret<-
  train(BadBuy ~ ., family='binomial', data=dfcTrain, method='lda',
        trControl = trainControl(method = "cv")) %>%
  predict(dfcTest, type='raw')%>%
  bind_cols(dfcTest, predictedClass=.)

resultsCaret%>%
  xtabs(~predictedClass+BadBuy, .)%>%
  confusionMatrix(positive='1')

## Confusion Matrix and Statistics
##
##              BadBuy
## predictedClass    0    1
##              0 1377  749
##              1  405  990
##
##              Accuracy : 0.6723
##              95% CI : (0.6565, 0.6878)
##              No Information Rate : 0.5061
##              P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.3428
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.5693
##              Specificity : 0.7727
##              Pos Pred Value : 0.7097
##              Neg Pred Value : 0.6477
##              Prevalence : 0.4939
##              Detection Rate : 0.2812
##              Detection Prevalence : 0.3962
##              Balanced Accuracy : 0.6710
##
##              'Positive' Class : 1
##
```

Q5)b)

```
set.seed(123)

ctrl <- trainControl(method="repeatedcv", number = 10, repeats = 3)
```

```
knnFit <- train(BadBuy ~ ., data = dfcTrain, method = "knn", trControl = ctrl
, preProcess = c("center","scale"), tuneLength = 20)
```

```
knnFit1<-knnFit%>%predict(dfcTest, type='raw')%>%
  bind_cols(dfcTest, predictedClass=.)
```

```
knnFit1%>%
  xtabs(~predictedClass+BadBuy, .)%>%
  confusionMatrix(positive='1')
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           BadBuy
```

```
## predictedClass  0    1
```

```
##           0 1322  815
```

```
##           1  460  924
```

```
##
```

```
##           Accuracy : 0.6379
```

```
##           95% CI : (0.6218, 0.6538)
```

```
##           No Information Rate : 0.5061
```

```
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.2739
```

```
##
```

```
##           Mcnemar's Test P-Value : < 2.2e-16
```

```
##
```

```
##           Sensitivity : 0.5313
```

```
##           Specificity : 0.7419
```

```
##           Pos Pred Value : 0.6676
```

```
##           Neg Pred Value : 0.6186
```

```
##           Prevalence : 0.4939
```

```
##           Detection Rate : 0.2624
```

```
##           Detection Prevalence : 0.3931
```

```
##           Balanced Accuracy : 0.6366
```

```
##
```

```
##           'Positive' Class : 1
```

```
##
```

```
#print(metrics.accuracy_score(y_test, y_pred))
```

```
knnFit
```

```
## k-Nearest Neighbors
```

```
##
```

```
## 6540 samples
```

```
## 9 predictor
```

```
## 2 classes: '0', '1'
```


```
##
```

```

## Pre-processing: centered (59), scaled (59)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 5885, 5886, 5887, 5887, 5885, 5886, ...
## Resampling results across tuning parameters:
##
##   k    Accuracy    Kappa
##   5  0.6040325  0.2079463
##   7  0.6045911  0.2090288
##   9  0.6063767  0.2125878
##  11  0.6119275  0.2236915
##  13  0.6106052  0.2210440
##  15  0.6148821  0.2295868
##  17  0.6189115  0.2376305
##  19  0.6196234  0.2390484
##  21  0.6184012  0.2365925
##  23  0.6181482  0.2360795
##  25  0.6207481  0.2412715
##  27  0.6199291  0.2396292
##  29  0.6204356  0.2406329
##  31  0.6232918  0.2463382
##  33  0.6246674  0.2490864
##  35  0.6300197  0.2597791
##  37  0.6291018  0.2579391
##  39  0.6310379  0.2618037
##  41  0.6335360  0.2668000
##  43  0.6343013  0.2683294
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 43.

plot(knnFit)

```



Q5)c)

```
set.seed(123)
```

```
#Set the grid for the Lambda values
```

```
lambdaValues <- 10^seq(-5, 2, length = 100)
```

```
fitLasso <- train(BadBuy ~ ., family='binomial', data=dfcTrain, method='glmnet',
  trControl=trainControl(method='cv', number=10), tuneGrid = expand.grid(alpha=1, lambda=lambdaValues))
```

```
#Variable importance complete table
```

```
varImp(fitLasso)$importance %>% # Add scale=FALSE inside VarImp if you don't want to scale
```

```
  rownames_to_column(var = "Variable") %>%
```

```
  mutate(Importance = scales::percent(Overall/100)) %>%
```

```
  arrange(desc(Overall)) %>%
```

```
  as_tibble()
```

```
## # A tibble: 59 x 3
```

```
##   Variable      Overall Importance
```


```
##   <chr>          <dbl> <chr>
```

```
## 1 WheelTypeNULL    100  100%
```

```
## 2 ColorOTHER       36.2  36%
```

```
## 3 SizeLARGETRUCK      26.5 26%
## 4 SizeCROSSOVER       24.2 24%
## 5 SizeLARGE            20.1 20%
## 6 MakeLINCOLN         19.7 20%
## 7 SizeLARGESUV        19.6 20%
## 8 MakeSUZUKI          19.4 19%
## 9 SizeSMALLTRUCK      18.8 19%
## 10 MakeMITSUBISHI     18.1 18%
## # ... with 49 more rows
```

```
#Variable importance plot with the most important variables
plot(varImp(fitLasso),top=25) # Add top = XX to change the number of visible
variables
```



```
#Optimum Lambda selected by the algorithm
fitLasso$bestTune$lambda # You can also run fitLasso$finalModel$lambdaOpt
```

```
## [1] 0.0003053856
```

```
#Not so useful but helps with understanding -See how variables are dropped as
lambda increases
```

```
plot(fitLasso$finalModel, xvar="lambda", label = TRUE)
```

```
#Not so useful but helps with understanding -See the coefficients from the fi
nal lasso model
```

```
#coef(fitLasso$finalModel, fitLasso$bestTune$lambda) # You can also use fit
Lasso$finalModel$lambdaOpt for optimum lambda
```



```

resultsLasso <-
  fitLasso %>%
  predict(dfctest, type='raw') %>%
  bind_cols(dfctest, predictedClass=.)

resultsLasso %>%
  xtabs(~predictedClass+BadBuy, .) %>%
  confusionMatrix(positive = '1')

## Confusion Matrix and Statistics
##
##               BadBuy
## predictedClass    0    1
##               0 1339  721
##               1  443 1018
##
##               Accuracy : 0.6694
##               95% CI : (0.6536, 0.6849)
##      No Information Rate : 0.5061
##      P-Value [Acc > NIR] : < 2e-16
##
##               Kappa : 0.3374
##
##  Mcnemar's Test P-Value : 4.7e-16
##
##               Sensitivity : 0.5854
##               Specificity : 0.7514
##               Pos Pred Value : 0.6968
##               Neg Pred Value : 0.6500
##               Prevalence : 0.4939
##               Detection Rate : 0.2891
##               Detection Prevalence : 0.4149
##               Balanced Accuracy : 0.6684
##
##               'Positive' Class : 1
##
fitLasso

## glmnet
##
## 6540 samples
##   9 predictor
##   2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 5885, 5886, 5887, 5887, 5885, 5886, ...
## Resampling results across tuning parameters:

```

##	lambda	Accuracy	Kappa
##	1.000000e-05	0.6746304	0.3490374
##	1.176812e-05	0.6746304	0.3490374
##	1.384886e-05	0.6746304	0.3490374
##	1.629751e-05	0.6746304	0.3490374
##	1.917910e-05	0.6746304	0.3490374
##	2.257020e-05	0.6746304	0.3490374
##	2.656088e-05	0.6746304	0.3490374
##	3.125716e-05	0.6746304	0.3490374
##	3.678380e-05	0.6746304	0.3490374
##	4.328761e-05	0.6746304	0.3490374
##	5.094138e-05	0.6746304	0.3490374
##	5.994843e-05	0.6746304	0.3490374
##	7.054802e-05	0.6746304	0.3490374
##	8.302176e-05	0.6746304	0.3490374
##	9.770100e-05	0.6746304	0.3490374
##	1.149757e-04	0.6746304	0.3490374
##	1.353048e-04	0.6746304	0.3490374
##	1.592283e-04	0.6746304	0.3490374
##	1.873817e-04	0.6746304	0.3490374
##	2.205131e-04	0.6746304	0.3490369
##	2.595024e-04	0.6743246	0.3484258
##	3.053856e-04	0.6747826	0.3493417
##	3.593814e-04	0.6744763	0.3487289
##	4.229243e-04	0.6741705	0.3481170
##	4.977024e-04	0.6738659	0.3475095
##	5.857021e-04	0.6721834	0.3441435
##	6.892612e-04	0.6724876	0.3447535
##	8.111308e-04	0.6718764	0.3435332
##	9.545485e-04	0.6717231	0.3432275
##	1.123324e-03	0.6724869	0.3447525
##	1.321941e-03	0.6721823	0.3441394
##	1.555676e-03	0.6723356	0.3444435
##	1.830738e-03	0.6717240	0.3432168
##	2.154435e-03	0.6721811	0.3441317
##	2.535364e-03	0.6717235	0.3432128
##	2.983647e-03	0.6712632	0.3422928
##	3.511192e-03	0.6705008	0.3407626
##	4.132012e-03	0.6692787	0.3383076
##	4.862602e-03	0.6688195	0.3373833
##	5.722368e-03	0.6666784	0.3330927
##	6.734151e-03	0.6669846	0.3336966
##	7.924829e-03	0.6683620	0.3364449
##	9.326033e-03	0.6691241	0.3379647
##	1.097499e-02	0.6680503	0.3358073
##	1.291550e-02	0.6689680	0.3376311
##	1.519911e-02	0.6688155	0.3373188
##	1.788650e-02	0.6675909	0.3348565
##	2.104904e-02	0.6682027	0.3360617

##	2.477076e-02	0.6668271	0.3333036
##	2.915053e-02	0.6665212	0.3326958
##	3.430469e-02	0.6660642	0.3317883
##	4.037017e-02	0.6666763	0.3330213
##	4.750810e-02	0.6666760	0.3330266
##	5.590810e-02	0.6666760	0.3330266
##	6.579332e-02	0.6562754	0.3119155
##	7.742637e-02	0.6559696	0.3113039
##	9.111628e-02	0.6406690	0.2804980
##	1.072267e-01	0.6209484	0.2408648
##	1.261857e-01	0.6212542	0.2414764
##	1.484968e-01	0.6212542	0.2414764
##	1.747528e-01	0.6212542	0.2414764
##	2.056512e-01	0.5009174	0.0000000
##	2.420128e-01	0.5009174	0.0000000
##	2.848036e-01	0.5009174	0.0000000
##	3.351603e-01	0.5009174	0.0000000
##	3.944206e-01	0.5009174	0.0000000
##	4.641589e-01	0.5009174	0.0000000
##	5.462277e-01	0.5009174	0.0000000
##	6.428073e-01	0.5009174	0.0000000
##	7.564633e-01	0.5009174	0.0000000
##	8.902151e-01	0.5009174	0.0000000
##	1.047616e+00	0.5009174	0.0000000
##	1.232847e+00	0.5009174	0.0000000
##	1.450829e+00	0.5009174	0.0000000
##	1.707353e+00	0.5009174	0.0000000
##	2.009233e+00	0.5009174	0.0000000
##	2.364489e+00	0.5009174	0.0000000
##	2.782559e+00	0.5009174	0.0000000
##	3.274549e+00	0.5009174	0.0000000
##	3.853529e+00	0.5009174	0.0000000
##	4.534879e+00	0.5009174	0.0000000
##	5.336699e+00	0.5009174	0.0000000
##	6.280291e+00	0.5009174	0.0000000
##	7.390722e+00	0.5009174	0.0000000
##	8.697490e+00	0.5009174	0.0000000
##	1.023531e+01	0.5009174	0.0000000
##	1.204504e+01	0.5009174	0.0000000
##	1.417474e+01	0.5009174	0.0000000
##	1.668101e+01	0.5009174	0.0000000
##	1.963041e+01	0.5009174	0.0000000
##	2.310130e+01	0.5009174	0.0000000
##	2.718588e+01	0.5009174	0.0000000
##	3.199267e+01	0.5009174	0.0000000
##	3.764936e+01	0.5009174	0.0000000
##	4.430621e+01	0.5009174	0.0000000
##	5.214008e+01	0.5009174	0.0000000
##	6.135907e+01	0.5009174	0.0000000
##	7.220809e+01	0.5009174	0.0000000

```
## 8.497534e+01 0.5009174 0.0000000
## 1.000000e+02 0.5009174 0.0000000
##
## Tuning parameter 'alpha' was held constant at a value of 1
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 1 and lambda = 0.00030538
56.
```

Q5)d)

```
set.seed(123)


#Set the grid for the Lambda values
lambdaValues <- 10^seq(-5, 2, length = 100)

fitRidge <- train(BadBuy ~ ., family='binomial', data=dfcTrain, method='glmnet',
  trControl=trainControl(method='cv', number=10), tuneGrid = expand.grid(alpha=0,
  lambda=lambdaValues))

#Variable importance complete table
varImp(fitRidge)$importance %>% # Add scale=FALSE inside VarImp if you don't want to scale
  rownames_to_column(var = "Variable") %>%
  mutate(Importance = scales::percent(Overall/100)) %>%
  arrange(desc(Overall)) %>%
  as_tibble()

## # A tibble: 59 x 3
##   Variable      Overall Importance
##   <chr>         <dbl> <chr>
## 1 WheelTypeNULL    100  100.00000%
## 2 MakeINFINITI     38.9  38.93548%
## 3 MakeLINCOLN      32.2  32.15789%
## 4 ColorOTHER       28.8  28.84262%
## 5 MakeOTHER        24.8  24.79872%
## 6 MakeOLDSMOBILE   21.5  21.50571%
## 7 MakeSCION        20.8  20.77209%
## 8 MakeSUZUKI       20.7  20.74687%
## 9 MakeMITSUBISHI   19.9  19.91029%
## 10 MakeGMC         17.5  17.50910%
## # ... with 49 more rows

#Variable importance plot with the most important variables
plot(varImp(fitRidge),top=25) # Add top = XX to change the number of visible variables
```



```
#Optimum lambda selected by the algorithm
fitRidge$bestTune$lambda # You can also run fitLasso$finalModel$lambdaOpt

## [1] 0.0559081

#Not so useful but helps with understanding -See how variables are dropped as
lambda increases
#plot(fitLasso$finalModel, xvar="lambda", label = TRUE)

#Not so useful but helps with understanding -See the coefficients from the fi
nal lasso model
#coef(fitLasso$finalModel, fitLasso$bestTune$lambda) # You can also use fit
Lasso$finalModel$lambdaOpt for optimum lambda

resultsRidge <-
  fitRidge %>%
  predict(dfctest, type='raw') %>%
  bind_cols(dfctest, predictedClass=.)

resultsRidge %>%
  xtabs(~predictedClass+BadBuy, .) %>%
  confusionMatrix(positive = '1')

## Confusion Matrix and Statistics
##
##               BadBuy
## predictedClass    0    1
```

```

##           0 1323  699
##           1  459 1040
##
##           Accuracy : 0.6711
##           95% CI : (0.6553, 0.6866)
##           No Information Rate : 0.5061
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.341
##
## Mcnemar's Test P-Value : 2.166e-12
##
##           Sensitivity : 0.5980
##           Specificity : 0.7424
##           Pos Pred Value : 0.6938
##           Neg Pred Value : 0.6543
##           Prevalence : 0.4939
##           Detection Rate : 0.2954
##           Detection Prevalence : 0.4257
##           Balanced Accuracy : 0.6702
##
##           'Positive' Class : 1
##
fitRidge

## glmnet
##
## 6540 samples
## 9 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 5885, 5886, 5887, 5887, 5885, 5886, ...
## Resampling results across tuning parameters:
##
##  lambda      Accuracy      Kappa
##  1.000000e-05 0.6720303 0.3438493
##  1.176812e-05 0.6720303 0.3438493
##  1.384886e-05 0.6720303 0.3438493
##  1.629751e-05 0.6720303 0.3438493
##  1.917910e-05 0.6720303 0.3438493
##  2.257020e-05 0.6720303 0.3438493
##  2.656088e-05 0.6720303 0.3438493
##  3.125716e-05 0.6720303 0.3438493
##  3.678380e-05 0.6720303 0.3438493
##  4.328761e-05 0.6720303 0.3438493
##  5.094138e-05 0.6720303 0.3438493
##  5.994843e-05 0.6720303 0.3438493

```

##	7.054802e-05	0.6720303	0.3438493
##	8.302176e-05	0.6720303	0.3438493
##	9.770100e-05	0.6720303	0.3438493
##	1.149757e-04	0.6720303	0.3438493
##	1.353048e-04	0.6720303	0.3438493
##	1.592283e-04	0.6720303	0.3438493
##	1.873817e-04	0.6720303	0.3438493
##	2.205131e-04	0.6720303	0.3438493
##	2.595024e-04	0.6720303	0.3438493
##	3.053856e-04	0.6720303	0.3438493
##	3.593814e-04	0.6720303	0.3438493
##	4.229243e-04	0.6720303	0.3438493
##	4.977024e-04	0.6720303	0.3438493
##	5.857021e-04	0.6720303	0.3438493
##	6.892612e-04	0.6720303	0.3438493
##	8.111308e-04	0.6720303	0.3438493
##	9.545485e-04	0.6720303	0.3438493
##	1.123324e-03	0.6720303	0.3438493
##	1.321941e-03	0.6720303	0.3438493
##	1.555676e-03	0.6720303	0.3438493
##	1.830738e-03	0.6720303	0.3438493
##	2.154435e-03	0.6720303	0.3438493
##	2.535364e-03	0.6720303	0.3438493
##	2.983647e-03	0.6720303	0.3438493
##	3.511192e-03	0.6720303	0.3438493
##	4.132012e-03	0.6720303	0.3438493
##	4.862602e-03	0.6720303	0.3438493
##	5.722368e-03	0.6720303	0.3438493
##	6.734151e-03	0.6720303	0.3438493
##	7.924829e-03	0.6720303	0.3438493
##	9.326033e-03	0.6720303	0.3438493
##	1.097499e-02	0.6720303	0.3438493
##	1.291550e-02	0.6720303	0.3438493
##	1.519911e-02	0.6720303	0.3438493
##	1.788650e-02	0.6718771	0.3435427
##	2.104904e-02	0.6714182	0.3426252
##	2.477076e-02	0.6724883	0.3447662
##	2.915053e-02	0.6726405	0.3450732
##	3.430469e-02	0.6724885	0.3447713
##	4.037017e-02	0.6724909	0.3447774
##	4.750810e-02	0.6721848	0.3441688
##	5.590810e-02	0.6734076	0.3466187
##	6.579332e-02	0.6727965	0.3454017
##	7.742637e-02	0.6711157	0.3420419
##	9.111628e-02	0.6692813	0.3383774
##	1.072267e-01	0.6688237	0.3374670
##	1.261857e-01	0.6705040	0.3408389
##	1.484968e-01	0.6703523	0.3405422
##	1.747528e-01	0.6689771	0.3377974
##	2.056512e-01	0.6686722	0.3371928

```

## 2.420128e-01 0.6685191 0.3368930
## 2.848036e-01 0.6674487 0.3347573
## 3.351603e-01 0.6669910 0.3338472
## 3.944206e-01 0.6674492 0.3347690
## 4.641589e-01 0.6674473 0.3347721
## 5.462277e-01 0.6671418 0.3341655
## 6.428073e-01 0.6691281 0.3381464
## 7.564633e-01 0.6686701 0.3372317
## 8.902151e-01 0.6677527 0.3354014
## 1.047616e+00 0.6674464 0.3347907
## 1.232847e+00 0.6668345 0.3335684
## 1.450829e+00 0.6669870 0.3338752
## 1.707353e+00 0.6668348 0.3335704
## 2.009233e+00 0.6662232 0.3323469
## 2.364489e+00 0.6659176 0.3317352
## 2.782559e+00 0.6662241 0.3323476
## 3.274549e+00 0.6665294 0.3329564
## 3.853529e+00 0.6663761 0.3326458
## 4.534879e+00 0.6659183 0.3317263
## 5.336699e+00 0.6665299 0.3329428
## 6.280291e+00 0.6666828 0.3332437
## 7.390722e+00 0.6672944 0.3344623
## 8.697490e+00 0.6679049 0.3356706
## 1.023531e+01 0.6663747 0.3325993
## 1.204504e+01 0.6659169 0.3316659
## 1.417474e+01 0.6686710 0.3371537
## 1.668101e+01 0.6683655 0.3365257
## 1.963041e+01 0.6680580 0.3358789
## 2.310130e+01 0.6668345 0.3334016
## 2.718588e+01 0.6639277 0.3275498
## 3.199267e+01 0.6610202 0.3216894
## 3.764936e+01 0.6596445 0.3188901
## 4.430621e+01 0.6542937 0.3081220
## 5.214008e+01 0.6504699 0.3004086
## 6.135907e+01 0.6415970 0.2825817
## 7.220809e+01 0.6307397 0.2607601
## 8.497534e+01 0.6272166 0.2536259
## 1.000000e+02 0.6188082 0.2367086
##
## Tuning parameter 'alpha' was held constant at a value of 0
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0 and lambda = 0.0559081.

```

Q5d)II)

```

# Elastic net with alpha = 0.5
# [*Naive elastic net - Read "Zou & Hastie (2005) Regularization and variable
selection via the elastic net" for details]

#Set the grid for the lambda values

```



```

lambdaValues <- 10^seq(-5, 2, length = 100)

set.seed(2020)


fitElastic <- train(BadBuy ~ ., family='binomial', data=dfcTrain, method='glm
net', trControl=trainControl(method='cv', number=10), tuneGrid=expand.grid(al
pha=0.5, lambda=lambdaValues))

#Variable importance complete table
varImp(fitElastic)$importance %>%      # Add scale=FALSE inside VarImp if you d
on't want to scale
  rownames_to_column(var = "Variable") %>%
  mutate(Importance = scales::percent(Overall/100)) %>%
  arrange(desc(Overall)) %>%
  as_tibble()

## # A tibble: 59 x 3
##   Variable      Overall Importance
##   <chr>         <dbl> <chr>
## 1 WheelTypeNULL    100  100%
## 2 ColorOTHER       35.9  36%
## 3 SizeLARGETRUCK   27.5  27%
## 4 SizeCROSSOVER    25.2  25%
## 5 SizeLARGESUV     20.8  21%
## 6 SizeLARGE        20.8  21%
## 7 MakeLINCOLN      20.1  20%
## 8 MakeSUZUKI       19.7  20%
## 9 SizeSMALLTRUCK   19.6  20%
## 10 SizeSMALLSUV    18.4  18%
## # ... with 49 more rows

#Variable importance plot with the most important variables
plot(varImp(fitElastic))      # Add top = XX to change the number of visible va
riables

```



```
#Optimum Lambda selected by the algorithm
fitElastic$bestTune$lambda # You can also run fitElastic$finalModel$LambdaOpt

## [1] 0.0003053856

#Not so useful but helps with understanding -See how variables are dropped as
#lambda increases
#plot(fitElastic$finalModel, xvar="lambda", label = TRUE)

#Not so useful but helps with understanding -See the coefficients from the final
#Elastic model
#coef(fitElastic$finalModel, fitElastic$bestTune$lambda) # You can also use
#fitElastic$finalModel$LambdaOpt for optimum lambda

resultsElastic <-
  fitElastic %>%
  predict(dfcTest, type='raw') %>%
  bind_cols(dfcTest, predictedClass=.)

resultsElastic %>%
  xtabs(~predictedClass+BadBuy, .) %>%
  confusionMatrix(positive = '1')

## Confusion Matrix and Statistics
##
##               BadBuy
```

```
## predictedClass    0    1
##                0 1342  720
##                1  440 1019
##
##                Accuracy : 0.6705
##                95% CI : (0.6547, 0.6861)
##      No Information Rate : 0.5061
##      P-Value [Acc > NIR] : < 2.2e-16
##
##                Kappa : 0.3397
##
##  McNemar's Test P-Value : 2.575e-16
##
##                Sensitivity : 0.5860
##                Specificity : 0.7531
##      Pos Pred Value : 0.6984
##      Neg Pred Value : 0.6508
##      Prevalence : 0.4939
##      Detection Rate : 0.2894
##      Detection Prevalence : 0.4144
##      Balanced Accuracy : 0.6695
##
##      'Positive' Class : 1
##
```

Q5)e)

```
set.seed(123)

resultsCaret<-
  train(BadBuy ~ ., family='binomial', data=dfcTrain, method='qda',
        trControl = trainControl(method = "cv")) %>%
  predict(dfcTest, type='raw')%>%
  bind_cols(dfcTest, predictedClass=.)

## Warning: model fit failed for Fold03: parameter=none Error in qda.default(
## x, grouping, ...) : rank deficiency in group 0
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainI
## nfo, :
## There were missing values in resampled performance measures.

resultsCaret%>%
  xtabs(~predictedClass+BadBuy, .)%>%
  confusionMatrix(positive='1')

## Confusion Matrix and Statistics
##
##                BadBuy
## predictedClass    0    1
##                0 1483  973
```

```
##          1  299  766
##
##          Accuracy : 0.6387
##          95% CI : (0.6226, 0.6546)
##    No Information Rate : 0.5061
##    P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 0.274
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##          Sensitivity : 0.4405
##          Specificity : 0.8322
##    Pos Pred Value : 0.7192
##    Neg Pred Value : 0.6038
##    Prevalence : 0.4939
##    Detection Rate : 0.2176
##    Detection Prevalence : 0.3025
##    Balanced Accuracy : 0.6363
##
##    'Positive' Class : 1
##
```